# Running a MongoDB Database in Kubernetes with StatefulSets

MongoDB is a widely used document-oriented NoSQL database that can be deployed in a containerized environment like Kubernetes. In this report, we will discuss the steps involved in running a MongoDB database in Kubernetes using StatefulSets in Qwiklab.

Steps:
Qwiklab provides a cloud-based platform to learn and practice various technical skills. To run a MongoDB database in Kubernetes using StatefulSets, we used the following steps:
Created a Kubernetes cluster using the Google Cloud Platform (GCP) Console and connected to it using the Cloud Shell.
Deployed a single-node MongoDB replica set using Kubernetes StatefulSets and Services.
Configured the MongoDB database for authentication and created a user with admin privileges.
Tested the MongoDB database by connecting to it using the mongo shell.

We were able to successfully deploy a MongoDB database in Kubernetes using StatefulSets and Services in Qwiklab. We configured the database for authentication and created a user with admin privileges. We also tested the database by connecting to it using the mongo shell.

Conclusion:
Running a MongoDB database in Kubernetes using StatefulSets is a useful skill for anyone working in a containerized environment. Qwiklab provides an excellent platform to learn and practice these skills, and we were able to successfully deploy a MongoDB database in Kubernetes using StatefulSets and Services.

Task 1: Set a compute zone

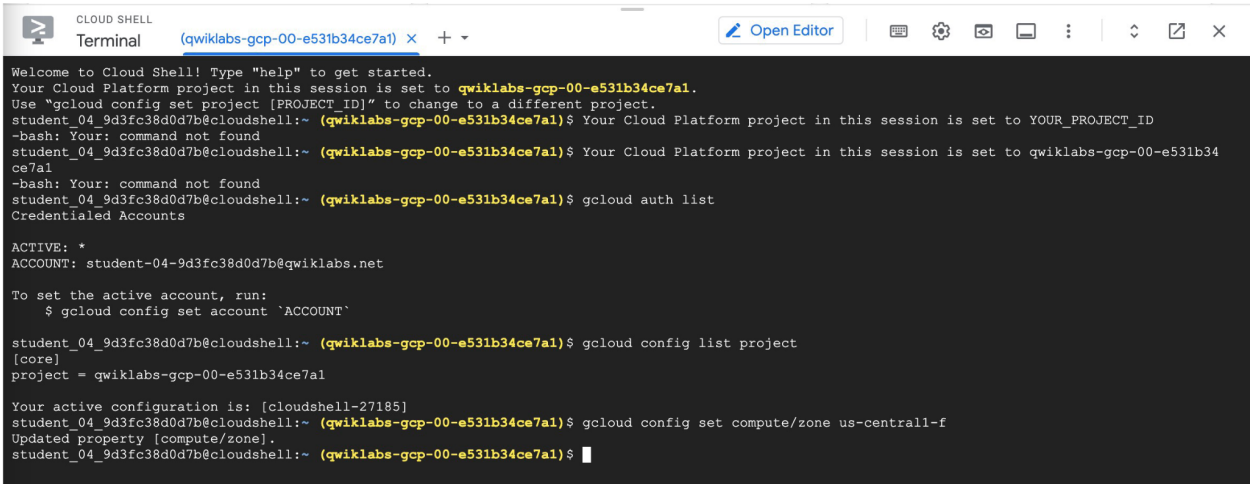


Task 2: Create a new cluster

## Task 3: Setting up

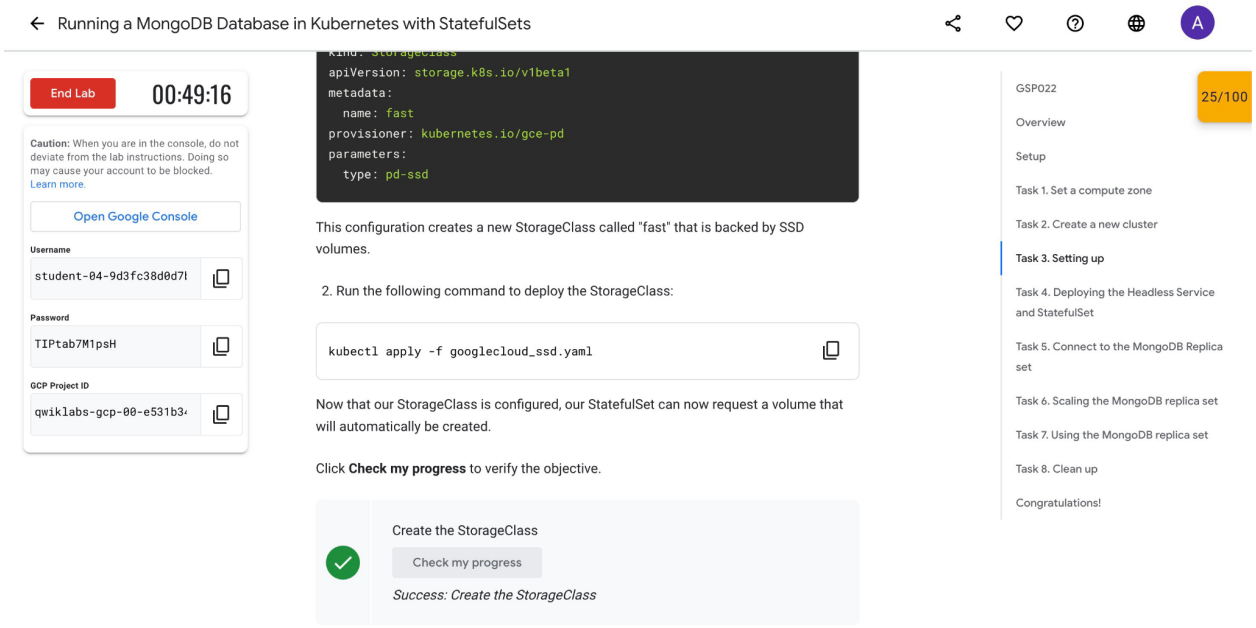


## Task 4: Deploying the headless Service and StatefulSet

## Task 5: Connect to the MongoDB Replica set

## Task 6: Scaling the MongoDB replica set

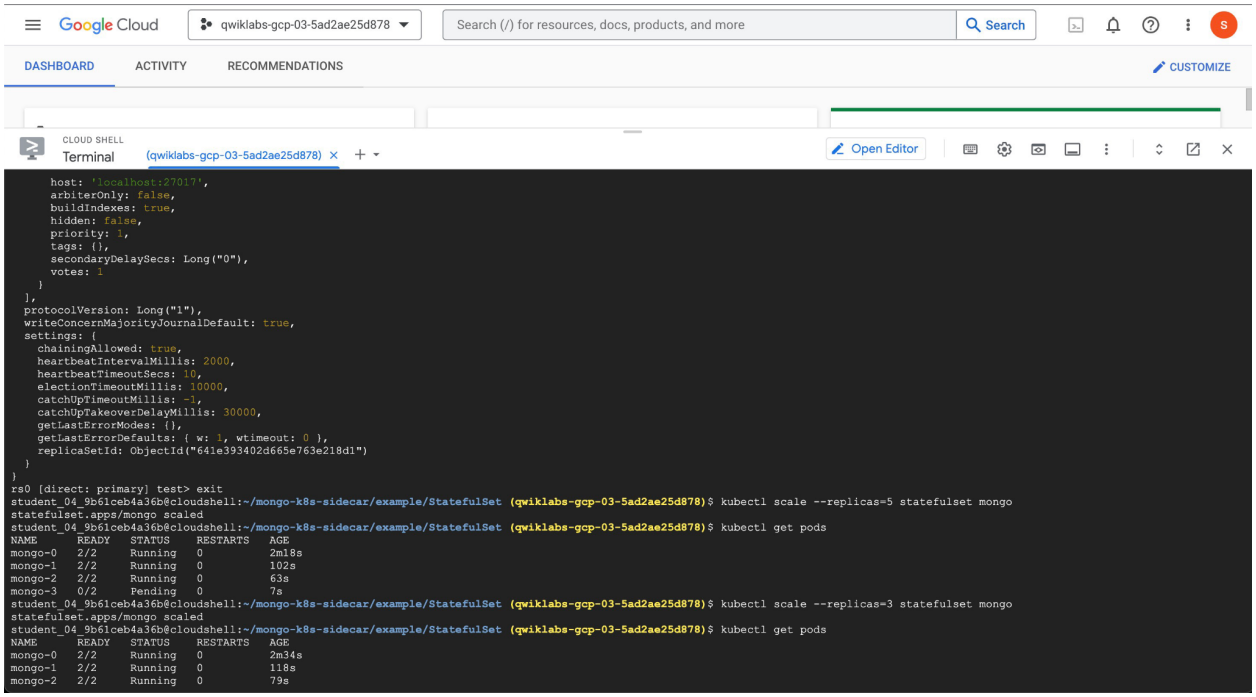


## Task 7: Using MongoDB replica set

Task 8: Clean Up
Congratulations:

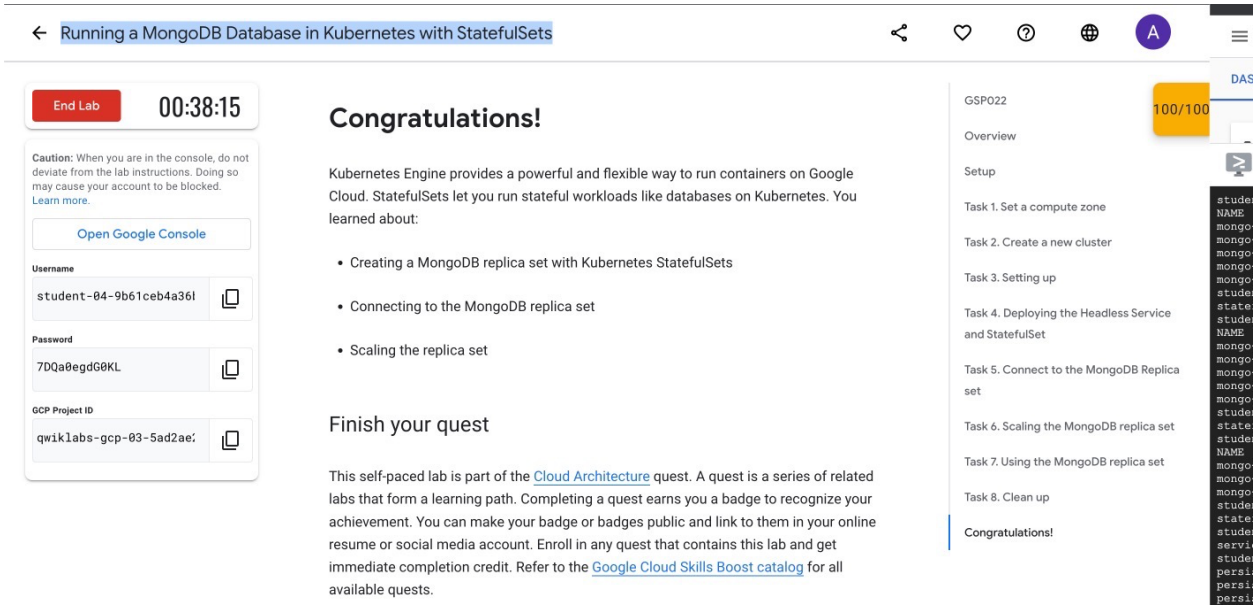


**For MongoDB, include your understanding of a headless service, StatefulSet, and why running a DB on a container is not a contradictory task**

In Kubernetes, a headless service is one that does not have a cluster IP allocated to it. Instead, it returns the IP addresses of the pods with which it is related. Headless services are utilized in StatefulSets to offer pods with stable network identities. A headless service is utilized in the case of a MongoDB deployment utilizing StatefulSets to ensure that each pod in the StatefulSet has a unique DNS entry that remains stable during the pod's lifespan.

StatefulSets are a sort of Kubernetes workload that assures the ordering and uniqueness of newly produced pods. They are helpful for stateful applications like databases because they ensure that each pod has a unique network identity that does not change even if the pod is rescheduled or removed.

Running a database in a container is not an incompatible task because containers offer a lightweight, portable, and scalable way to package and deploy applications. Databases, such as MongoDB, may be packaged as container images and deployed in a containerized environment such as Kubernetes. Containerization has various advantages, including isolation, repeatability, and scalability. Furthermore, operating a database in a containerized environment like Kubernetes enables more effective resource use and simplifies database infrastructure management as code.

# Ingesting New Datasets into BigQuery

Task 1. Create a new dataset to store tables
Task 2. Ingest a new Dataset from a CSV
Task 3. Ingest data from Cloud Storage
Task 4. Ingest a new dataset from a Google Spreadsheet

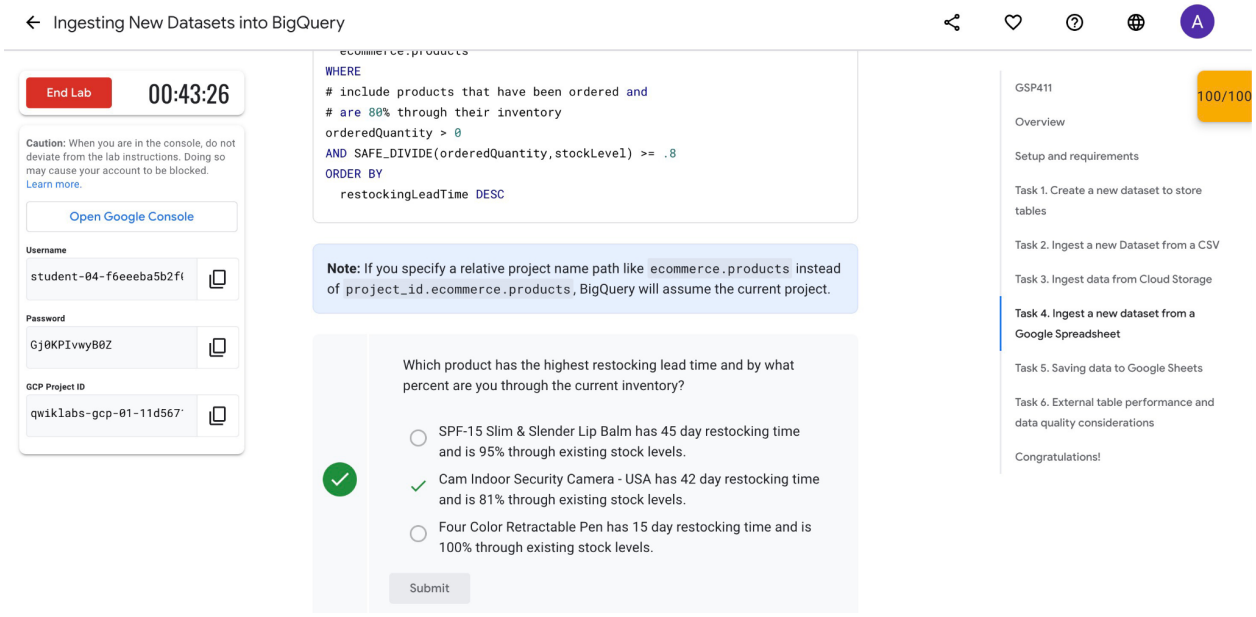


Task 5. Saving data to Google Sheets

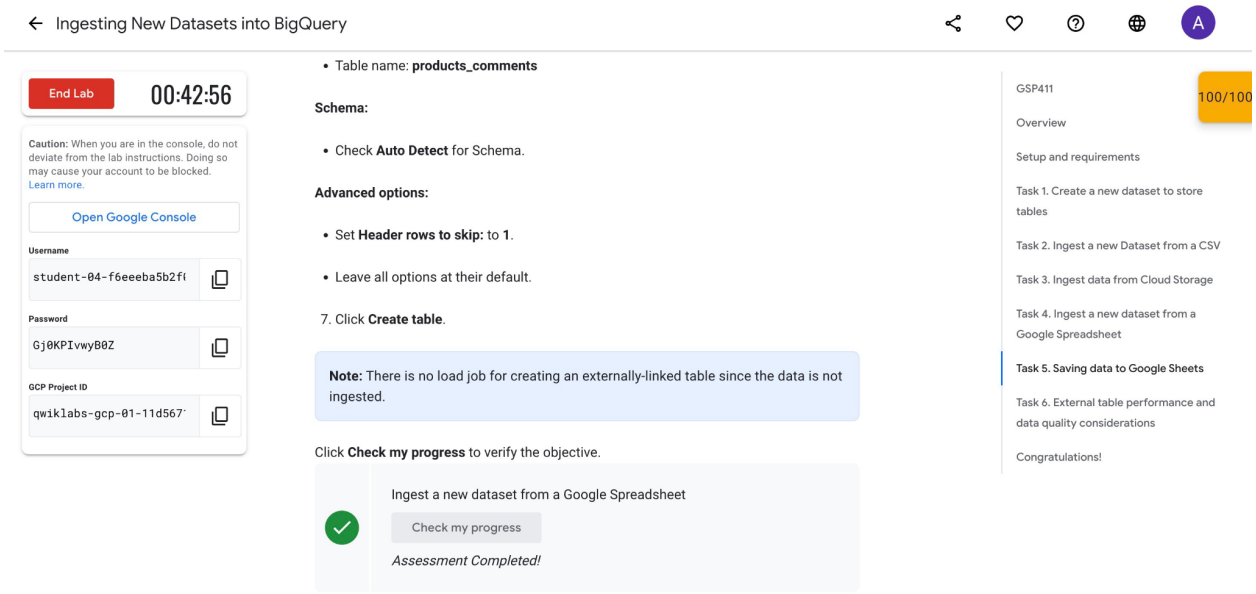

Task 6. External table performance and data quality considerations
Congratulations!

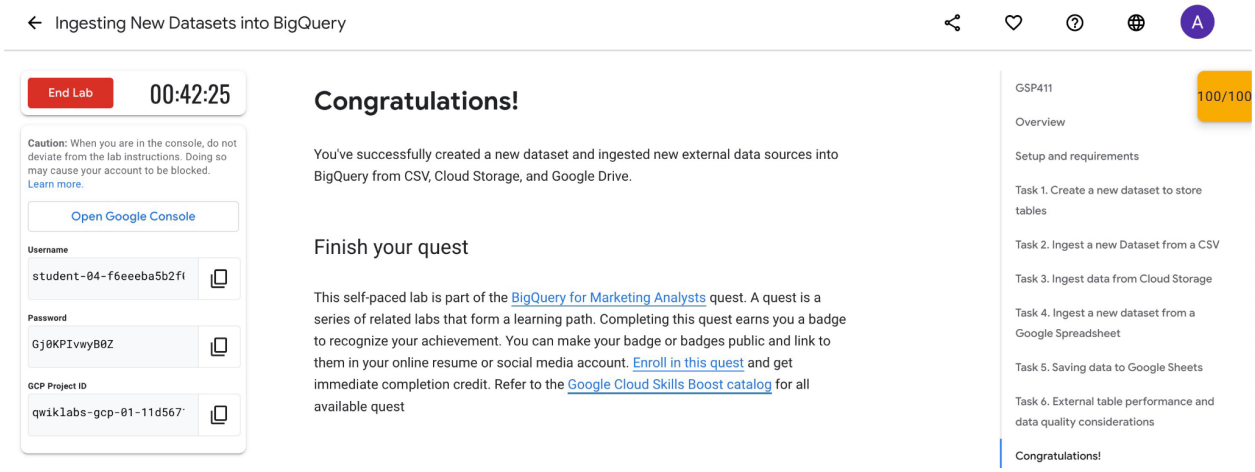


Qwiklabs is a cloud-based learning platform that provides hands-on training and practice labs on various cloud computing technologies. One of the labs offered by Qwiklabs is "Ingesting New Datasets into BigQuery," which is aimed at teaching users how to load new datasets into Google's BigQuery data warehouse.

In this lab, users are provided with a dataset in CSV format and are guided through the process of uploading the data into a new table in BigQuery. The lab includes step-by-step instructions on how to create a new dataset, create a new table, and load the data into the table using the Google Cloud Console.

The lab also covers various options for loading data, including loading data from a local file or a Google Cloud Storage bucket. Additionally, users learn how to configure the load job to handle various data formats, including CSV, JSON, and Avro.

**Pros and Cons of Big Query**
Pros:

1. BigQuery is built to handle petabyte-scale data and can manage enormous data intake and querying.
2. BigQuery is geared for speed, and queries may be executed in seconds, even on very huge datasets.
3. BigQuery is a pay-as-you-go service, which means you only pay for the data you consume and the processing power you require. Also, storage expenses are inexpensive, and there are no upfront expenditures.
4. BigQuery works nicely with other Google Cloud services as well as external tools like Tableau and Data Studio.
5. BigQuery has robust security features like encryption at rest and in transit, access control, and audit tracking.

Cons:

1. BigQuery is user-friendly, but there is a learning curve to fully understand its capabilities and how to use it effectively.
2. While BigQuery supports the majority of SQL functions, some are still unavailable, which might be a constraint for specific use cases.
3. Because moving data into and out of BigQuery can be expensive, it is critical to design your data storage and access strategy properly.
4. While BigQuery is generally inexpensive, conducting complicated searches or queries on extremely big datasets might incur considerable fees.
5. Access control is limited to the project- and dataset-level permissions in BigQuery, which may not be sufficient for some businesses with sophisticated data access requirements.